

# Depth Cameras, Demystified

Tutorial at the  
2018 Summer School on

Multimodal Interaction  
in Augmented and Virtual Reality

by Jun.-Prof. Dr. Florian Echtler  
<[florian.echtler@uni-weimar.de](mailto:florian.echtler@uni-weimar.de)>

# About me


- [@floemuc](https://floe.butterbrot.org/) & <https://floe.butterbrot.org/>
- „Postdoctoral Hacker and Research Hobbyist“
- Co-Author of `libfreenect{2}` open-source drivers for Kinect 1 & 2 depth cameras
- <https://github.com/OpenKinect/libfreenect2>
- <https://doi.org/10.5281/zenodo.50641>

# Plan for today

- 09:00 ~ 10:00 – Tutorial: depth camera basics
- 10:00 ~ 10:15 – short break
- 10:15 ~ 11:30 – Hands-on depth cam hacking  
(Build Your Own RANSAC)

# What are depth cameras?

- „Regular“ camera: color values for each pixel
- Depth camera: distance values for each pixel (usually visualized with color map)




# Classes of DCs

- Geometry-based
  - Stereo vision
  - Structured light
- Time-of-Flight/ToF

# Geometry-based DCs

- Fundamental principle:
  - Create two „views“ of scene
  - Match scene points between views
  - Determine 2 angles for each scene point
  - Trigonometry happens
  - Receive distance

# Geometry-based DCs



# Geometry: Stereo Matching

- Method 1: two images of scene
  - *stereo matching* of corresponding pixels
  - ideally only needed on horizontal *scanline*
- Examples: Occipital Structure Sensor, Intel Realsense D4xx



# Geometry: Stereo Matching




Image Source (FU): <https://www.mtbs3d.com/phpbb/viewtopic.php?f=138&t=18055>

# Geometry: Stereo Matching

- Advantages:
  - Close to human vision system
  - Can work outdoors, even in sunlight
  - Only uses plain cameras (\* with sync)

# Geometry: Stereo Matching

- Drawbacks:
  - Problems with featureless areas (white wall)
  - Mitigation: additional IR pattern projector



# Geometry-based DCs

- Method 2: one image of scene + known lighting
  - identify pixel correspondences via IR pattern
  - Replaces one camera with light source
  - 2 sub-methods (down the rabbit hole... :-)

# Geometry: Speckle Pattern

- Method 2.1: *speckle* pattern = random dots
  - Random, but previously known pattern
  - „Patches“ of pattern can be matched → depth res. lower than image res.
  - Example: Kinect v1



# Geometry: Stripe Pattern

- Method 2.2: Gray code = alternating stripes
  - Encodes binary ID for each pixel
  - Requires high frame rate or static scene (why?)
  - Requires IR projector
  - Example: Realsense SR300 (~ 300 FPS)





Image Source (FU): <http://www.sci.utah.edu/~gerig/.../CS6320-CV-S2012-StructuredLight-II.pdf>

# Time-of-Flight/ToF DCs

- Fundamental principle:
  - Emit (infrared) flash
  - Measure time until reflected light arrives
  - (very simple) Math happens
  - Receive distance
- Example: Kinect v2



# Time-of-Flight DCs





# Time-of-Flight DCs

- Problem:  $c \sim 300\,000\,000\text{ m/s}$  :-O ( $\sim 671\text{ mil. mph}$  ;-)
- Common ToF depth resolution  $\sim 1\text{ mm}$ 
  - time difference  $\sim 3 \times 10^{-12}\text{ s}$  (*3 picoseconds!*)
  - would require counter at  $\sim 300\text{ GHz}$  per pixel

HOW?


# Time-of-Flight DCs

- Solution: measure *phase* difference (not time)
  - Create clock signal with freq.  $f$  which modulates ...
    - a) IR emitter („flash“)
    - b) sensitivity of light sensor
  - Result is phase difference
  - Calculation with  $c$  and  $f$  gives distance

# Time-of-Flight DCs

- Requires specialized image sensor
  - Two “accumulators” per pixel (not just one)
  - Clock input switches between acc. A and B
- Kinect v2: result is three measurements per pixel at phase shift  $0^\circ$ ,  $120^\circ$  and  $240^\circ$

# Time-of-Flight DCs



## Differential Pixel

- $(A+B)$  gives the ambient (room) lighting ('common mode') – 'normal' grey scale image
- $(A-B)$  gives phase (depth) information after an arctan calculation – depth image
- $\sqrt{\Sigma(A-B)^2}$  is the 'Active' image – A grey scale image independent of ambient lighting

Image Source (FU): [https://www.hotchips.org/wp-content/uploads/hc\\_archives/hc25/Hc25.10-SoC1-epub/Hc25.26.121-fixed-%20XB1%2020130826gnn.pdf](https://www.hotchips.org/wp-content/uploads/hc_archives/hc25/Hc25.10-SoC1-epub/Hc25.26.121-fixed-%20XB1%2020130826gnn.pdf)

# Time-of-Flight DCs

- Tradeoff between precision and range
  - High frequency → better precision, but quicker “wraparound”
  - E.g. Kinect v2 @ 80 MHz: wavelength = 3.75 m
  - Range from 0 to 3.75 m maps to depth resolution of sensor (11 bit for Kv2)

# Time-of-Flight DCs

- Solution: switch between multiple frequencies
  - Kinect v2: 16 MHz, 80 MHz, 120 MHz
  - Resolve ambiguity using second/third freq.




Image Source (FU): [https://www.hotchips.org/wp-content/uploads/hc\\_archives/hc25/HC25.10-SoC1-epub/HC25.26.121-fixed-%20XB1%2020130826gmn.pdf](https://www.hotchips.org/wp-content/uploads/hc_archives/hc25/HC25.10-SoC1-epub/HC25.26.121-fixed-%20XB1%2020130826gmn.pdf)

# Depth-Color Alignment


- most depth cameras also have a “plain” color cam
- Problem: how to find corresponding color value for each depth pixel (or vice versa)?
- Requires intrinsic and extrinsic camera parameters (?)
  - Intrinsic: field of view, distortion, focal length, ...
  - Extrinsic: translation, rotation w.r.t. origin

# Depth-Color Alignment

- Alignment process:
  - “Deproject” depth pixels into metric space  
(needs *intrinsic* parameters of *depth cam*)
  - Translate depth pixels to color camera origin  
(needs *extrinsic* parameters of *depth cam*)
  - Cast ray in metric space through color pixel  
(needs *intrinsic* parameters of *color cam*)
  - Find (closest) intersection





# Depth-Color Alignment




# Point Clouds

- Each depth pixel  $\rightarrow$  3D point  $(x, y, z)$
- Plus color  $\rightarrow (x, y, z, r, g, b)$



# Question

- Why the shadow on the left side of the body?



# Summary

- One size does not fit all (as always)
  - Lighting? (e.g. sunlight – very hard)
  - Texture? (e.g. flat featureless walls)
  - Absorption? (e.g. some fabrics)
  - Reflections? (nearly impossible ATM)

# RANSAC

- RANSAC = RANdom SAmple Consensus
  - Goal: fit a model to noisy data
  - Model: e.g. plane, line, ...
  - Data: e.g. point cloud

# RANSAC

- Step 1: randomly choose minimum # of data points required for model (e.g. 2 for line)
- Step 2: count data points which are *inliers* (i.e. “close enough” to the model)
- Too few inliers → step 1
- Otherwise: refine model using inliers (optional)




Image Source (CC): <https://commons.wikimedia.org/w/index.php?curid=37017886>

# Demo Time!

- See <https://github.com/floe/surface-streams>
- and <https://www.youtube.com/watch?v=Qe1BROtGyzI>



# Your turn!

- Pick a depth camera from the pile
- Install the respective SDK
- Adapt the sample code to measure average deviation from plane
- Find a plane
- Profit!